

In []:

In []:

FIAT500_VEHICLE SELECTION USING *ELASTIC NET*

In []:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df=pd.read_csv(r"/content/fiat500_VehicleSelection_Dataset (1).csv")
df
```

Out[]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|------|------|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0 | 1 | lounge | 51 | 882 | 25000 | 1 | 44.907242 | 8.611560 | 8900 |
| 1 | 2 | pop | 51 | 1186 | 32500 | 1 | 45.666359 | 12.241890 | 8800 |
| 2 | 3 | sport | 74 | 4658 | 142228 | 1 | 45.503300 | 11.417840 | 4200 |
| 3 | 4 | lounge | 51 | 2739 | 160000 | 1 | 40.633171 | 17.634609 | 6000 |
| 4 | 5 | pop | 73 | 3074 | 106880 | 1 | 41.903221 | 12.495650 | 5700 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1533 | 1534 | sport | 51 | 3712 | 115280 | 1 | 45.069679 | 7.704920 | 7900 |
| 1534 | 1535 | lounge | 74 | 3835 | 112000 | 1 | 45.845692 | 8.666870 | 4200 |
| 1535 | 1536 | pop | 51 | 2223 | 60457 | 1 | 45.481541 | 9.413480 | 6000 |
| 1536 | 1537 | lounge | 51 | 2557 | 80750 | 1 | 45.000702 | 7.682270 | 7900 |
| 1537 | 1538 | pop | 51 | 1766 | 54276 | 1 | 40.323410 | 17.568270 | 6000 |

1538 rows × 9 columns



In []:

```
df.head(10)
```

Out[]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|---|----|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0 | 1 | lounge | 51 | 882 | 25000 | 1 | 44.907242 | 8.611560 | 8900 |
| 1 | 2 | pop | 51 | 1186 | 32500 | 1 | 45.666359 | 12.241890 | 8800 |
| 2 | 3 | sport | 74 | 4658 | 142228 | 1 | 45.503300 | 11.417840 | 4200 |
| 3 | 4 | lounge | 51 | 2739 | 160000 | 1 | 40.633171 | 17.634609 | 6000 |
| 4 | 5 | pop | 73 | 3074 | 106880 | 1 | 41.903221 | 12.495650 | 5700 |
| 5 | 6 | pop | 74 | 3623 | 70225 | 1 | 45.000702 | 7.682270 | 7900 |
| 6 | 7 | lounge | 51 | 731 | 11600 | 1 | 44.907242 | 8.611560 | 10750 |
| 7 | 8 | lounge | 51 | 1521 | 49076 | 1 | 41.903221 | 12.495650 | 9190 |
| 8 | 9 | sport | 73 | 4049 | 76000 | 1 | 45.548000 | 11.549470 | 5600 |
| 9 | 10 | sport | 51 | 3653 | 89000 | 1 | 45.438301 | 10.991700 | 6000 |



```
In [ ]: df.tail()
```

```
Out[ ]:
```

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|-------------|------|--------|--------------|-------------|--------|-----------------|-----------|----------|---------|
| 1533 | 1534 | sport | 51 | 3712 | 115280 | 1 | 45.069679 | 7.70492 | 5100000 |
| 1534 | 1535 | lounge | 74 | 3835 | 112000 | 1 | 45.845692 | 8.66687 | 4000000 |
| 1535 | 1536 | pop | 51 | 2223 | 60457 | 1 | 45.481541 | 9.41348 | 7000000 |
| 1536 | 1537 | lounge | 51 | 2557 | 80750 | 1 | 45.000702 | 7.68227 | 5000000 |
| 1537 | 1538 | pop | 51 | 1766 | 54276 | 1 | 40.323410 | 17.56827 | 7000000 |

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   int64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   int64
3   age_in_days           1538 non-null   int64
4   km                    1538 non-null   int64
5   previous_owners       1538 non-null   int64
6   lat                   1538 non-null   float64
7   lon                   1538 non-null   float64
8   price                 1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

```
In [ ]: df.describe()
```

```
Out[ ]:
```

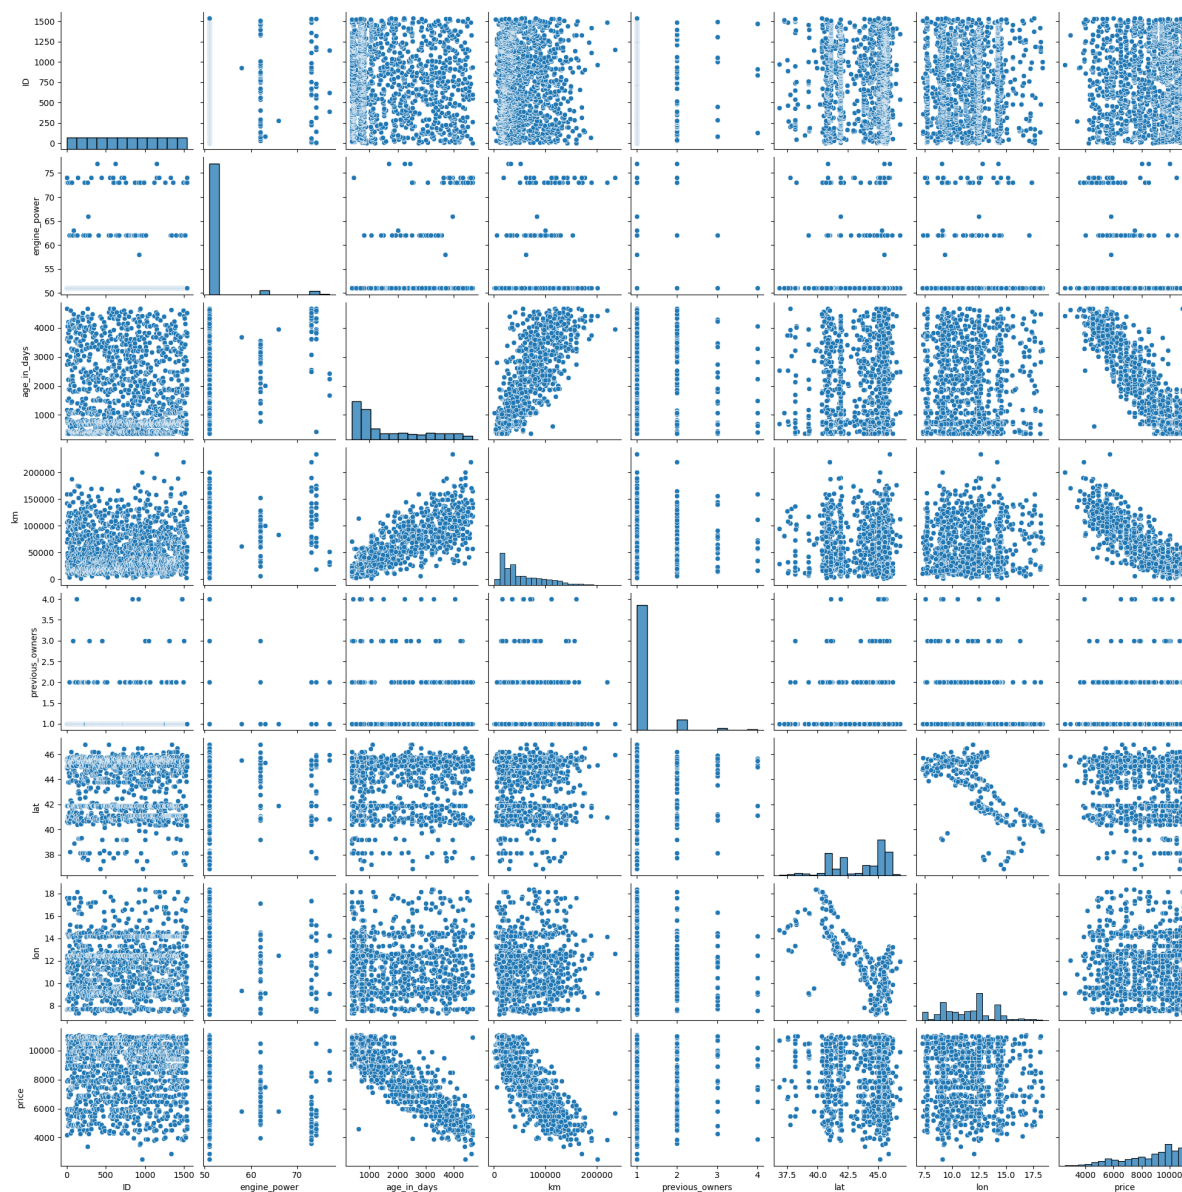
| | ID | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|--------------|-------------|--------------|-------------|---------------|-----------------|-------------|-------------|-------------|
| count | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 |
| mean | 769.500000 | 51.904421 | 1650.980494 | 53396.011704 | 1.123537 | 43.541361 | 10.123456 | 5000000.0 |
| std | 444.126671 | 3.988023 | 1289.522278 | 40046.830723 | 0.416423 | 2.133518 | 1.234567 | 1000000.0 |
| min | 1.000000 | 51.000000 | 366.000000 | 1232.000000 | 1.000000 | 36.855839 | 7.682227 | 4000000.0 |
| 25% | 385.250000 | 51.000000 | 670.000000 | 20006.250000 | 1.000000 | 41.802990 | 8.666870 | 5000000.0 |
| 50% | 769.500000 | 51.000000 | 1035.000000 | 39031.000000 | 1.000000 | 44.394096 | 10.123456 | 5000000.0 |
| 75% | 1153.750000 | 51.000000 | 2616.000000 | 79667.750000 | 1.000000 | 45.467960 | 17.568270 | 7000000.0 |
| max | 1538.000000 | 77.000000 | 4658.000000 | 235000.000000 | 4.000000 | 46.795612 | 17.568270 | 7000000.0 |

```
In [ ]: df.columns
```

```
Out[ ]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
            'lat', 'lon', 'price'],
            dtype='object')
```


```
In [ ]: #EDA
sns.pairplot(df)
```

Out[]: <seaborn.axisgrid.PairGrid at 0x7fed1e05cb20>




In []: `sns.displot(['Price'])`

Out[]: <seaborn.axisgrid.FacetGrid at 0x7fed19247070>




```
In [ ]: sns.displot(df['price'])
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7fed19246b30>
```




```
In [ ]: sns.displot(df['engine_power'])
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7fed16a63bb0>
```




```
In [ ]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                'lat', 'lon', 'price']]  
sns.heatmap(fiatdf.corr())
```

```
Out[ ]: <Axes: >
```



```
In [ ]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                'lat', 'lon']]  
sns.heatmap(fiatdf.corr())#without price
```

```
Out[ ]: <Axes: >
```



```
In [ ]: X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
               'lat', 'lon']]
        y=df['price']
```

```
In [ ]: from sklearn.model_selection import train_test_split
        X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)
        from sklearn.linear_model import LinearRegression
        regr=LinearRegression()
        regr.fit(X_train,y_train)
        print(regr.intercept_)
```

8971.195683500588


```
In [ ]: coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])
        coeff_df
```

```
Out[ ]:
```

| | coefficient |
|-----------------|-------------|
| ID | -0.046704 |
| engine_power | 11.646408 |
| age_in_days | -0.898018 |
| km | -0.017232 |
| previous_owners | 26.400886 |
| lat | 32.189709 |
| lon | 0.161073 |



```
In [ ]: predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

```
Out[ ]: <matplotlib.collections.PathCollection at 0x7fed138bd090>
```




```
In [ ]: sns.displot((y_test,predictions),bins=50)#without semicolon
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7fed198fbaf0>
```



```
In [ ]: sns.displot((y_test,predictions),bins=50);#with semicolon
```



```
In [ ]: from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('MAE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

MAE: 593.0876179519936

MSE: 551442.6799691812

MAE: 742.5918663500033

```
In [ ]: #accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.8597136704308864


```
In [ ]: df.fillna(method='ffill',inplace=True)
```

```
In [ ]: x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```

```
In [ ]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

```
Out[ ]: ▼ LinearRegression
LinearRegression()
```

```
In [ ]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [ ]: #elasticnet
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

[25.89689696]

[10640.73996329]

Mean Squared Error on test set 2692062172.9534926